

Image Recognition with IBM Cloud Visual Recognition

Phase 4: Development Part 2

OVERVIEW:

The following actions must be taken in order to construct an image recognition system that integrates AI-generated captions with IBM Cloud Visual Recognition:

Become an IBM Cloud user: Make sure you have an IBM Cloud account if you don't have. Establish a Visual Recognition Service: Launch an IBM Watson Visual Recognition service instance on IBM Cloud. You can use this service to analyze photos and find items in them. Obtain API Credentials: Once the Visual Recognition service has been created, get your API Key and URL. To authenticate and gain access to the service, you will require these credentials. Create the Image Classification Component: To connect to the Visual Recognition service, use Python and the IBM Watson SDK. This element examines the photos and provides a list of things it has identified.

Integrate Natural Language Generation (NLG): To generate captions for the objects that have been identified, utilize a library for natural language generation. For every object, you can create a coherent statement that describes it. Using the generated captions and the results of the image categorization, show the captions on a web page or user interface.

CODE FOR IMAGE RECOGNITION:

```
1. Install Dependencies and Setup

In [1]: !pip install tensorflow tensorflow-gpu opencv-python matplotlib

Requirement already satisfied: tensorflow in c:\users\suban_sri.rajabagal\anaconda3\lib\site-packages (2.14.0)
Collecting tensorflow-gpu
  Downloading tensorflow-gpu-2.12.0.tar.gz (2.6 kB)
  Preparing metadata (setup.py): started
  Preparing metadata (setup.py): finished with status 'done'
Collecting opencv-python
  Downloading opencv_python-4.8.1.78-cp37-abi3-win_amd64.whl (38.1 MB)
  ----- 38.1/38.1 MB 567.9 kB/s eta 0:00:00
Requirement already satisfied: matplotlib in c:\users\suban_sri.rajabagal\anaconda3\lib\site-packages (3.7.0)
Requirement already satisfied: tensorflow-intel==2.14.0 in c:\users\suban_sri.rajabagal\anaconda3\lib\site-packages (from tensorflow) (2.14.0)
Requirement already satisfied: h5py>=2.9.0 in c:\users\suban_sri.rajabagal\anaconda3\lib\site-packages (from tensorflow-intel==2.14.0->tensorflow) (3.7.0)
Requirement already satisfied: absl-py>=1.0.0 in c:\users\suban_sri.rajabagal\anaconda3\lib\site-packages (from tensorflow-intel==2.14.0->tensorflow) (2.0.0)
Requirement already satisfied: protobuf!=4.21.0,!=4.21.1,!=4.21.2,!=4.21.3,!=4.21.4,!=4.21.5,<5.0.0dev,>=3.20.3 in c:\users\suban_sri.rajabagal\anaconda3\lib\site-packages (from tensorflow-intel==2.14.0->tensorflow) (4.24.4)
Requirement already satisfied: astunparse>=1.6.0 in c:\users\suban_sri.rajabagal\anaconda3\lib\site-packages (from tensorflow-intel==2.14.0->tensorflow) (1.6.3)

In [2]: !pip list

Package                                Version
-----
absl-py                                2.0.0
alabaster                               0.7.12
anaconda-client                         1.11.2
anaconda-navigator                     2.4.2
anaconda-project                       0.11.1
anyio                                   3.5.0
appdirs                                1.4.4
argon2-cffi                            21.3.0
argon2-cffi-bindings                   21.2.0
```

```
In [3]: import tensorflow as tf
import os
```

```
decorator          5.1.1
defusedxml          0.7.1
diff-match-patch    20200713
dill                0.3.6
distributed         2022.7.0
docstring-to-markdown 0.11
docutils            0.18.1
entrypoints         0.4
et-xmlfile          1.1.0
executing           0.8.3
fastjsonschema      2.16.2
filelock            3.9.0
flake8              6.0.0
Flask               2.2.2
flatbuffers         23.5.26
flit_core           3.6.0
fonttools           4.25.0
fsspec              2022.11.0
```

```
In [ ]: # Avoid OOM errors by setting GPU Memory Consumption Growth
gpus = tf.config.experimental.list_physical_devices('GPU')
for gpu in gpus:
    tf.config.experimental.set_memory_growth(gpu, True)
```

```
In [ ]: tf.config.list_physical_devices('GPU')
```

2. Remove dodgy images

```
In [ ]: import cv2
import imghdr
```

```
In [ ]: data_dir = 'data'
```

```
In [ ]: image_exts = ['.jpeg', '.jpg', '.png', '.bmp', '.gif', '.tiff', '.webp']
```

```
In [ ]: for image_class in os.listdir(data_dir):
    for image in os.listdir(os.path.join(data_dir, image_class)):
        image_path = os.path.join(data_dir, image_class, image)
        try:
            img = cv2.imread(image_path)
            tip = imghdr.what(image_path)
            if tip not in image_exts:
                print('Image not in ext list {}'.format(image_path))
                os.remove(image_path)
        except Exception as e:
            print('Issue with image {}'.format(image_path))
            # os.remove(image_path)
```

3. Load Data

```
In [8]: import numpy as np
from matplotlib import pyplot as plt
```


```
In [9]: data = tf.keras.utils.image_dataset_from_directory('data')
```

Found 305 files belonging to 2 classes.

```
In [10]: data_iterator = data.as_numpy_iterator()
```

```
In [11]: batch = data_iterator.next()
```

```
In [12]: fig, ax = plt.subplots(ncols=4, figsize=(20,20))
for idx, img in enumerate(batch[0][:4]):
    ax[idx].imshow(img.astype(int))
    ax[idx].title.set_text(batch[1][idx])
```



```
In [ ]: data.as_numpy_iterator().next()
```

5. Split Data

```
In [15]: train_size = int(len(data)*.7)
val_size = int(len(data)*.2)
test_size = int(len(data)*.1)
```

```
In [16]: train_size
```

```
Out[16]: 7
```

```
In [17]: train = data.take(train_size)
val = data.skip(train_size).take(val_size)
test = data.skip(train_size+val_size).take(test_size)
```

6. Build Deep Learning Model

```
In [18]: train
```

```
Out[18]: <TakeDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None, 256, 256, 3), dtype=tf.int32, name=None))>
```

```
In [ ]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
```

```
In [ ]: model = Sequential()
```

```
In [ ]: model.add(Conv2D(16, (3,3), 1, activation='relu', input_shape=(256,256,3)))
model.add(MaxPooling2D())
model.add(Conv2D(32, (3,3), 1, activation='relu'))
model.add(MaxPooling2D())
model.add(Conv2D(16, (3,3), 1, activation='relu'))
model.add(MaxPooling2D())
model.add(Flatten())
```

```
Out[18]: <TakeDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None, 256, 256, 3), dtype=tf.int32, name=None))>
```

```
In [ ]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
```

```
In [ ]: model = Sequential()
```

```
In [ ]: model.add(Conv2D(16, (3,3), 1, activation='relu', input_shape=(256,256,3)))
model.add(MaxPooling2D())
model.add(Conv2D(32, (3,3), 1, activation='relu'))
model.add(MaxPooling2D())
model.add(Conv2D(16, (3,3), 1, activation='relu'))
model.add(MaxPooling2D())
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
```

```
In [ ]: model.compile('adam', loss=tf.losses.BinaryCrossentropy(), metrics=['accuracy'])
```

```
In [ ]: model.summary()
```

7. Train

```
In [ ]: logdir='logs'
```

```
In [ ]: tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=logdir)
```

```
In [ ]: hist = model.fit(train, epochs=20, validation_data=val, callbacks=[tensorboard_callback])
```

8. Plot Performance

```
from IPython.display import Image
```

8. Plot Performance

```
In [ ]: fig = plt.figure()
plt.plot(hist.history['loss'], color='teal', label='loss')
plt.plot(hist.history['val_loss'], color='orange', label='val_loss')
fig.suptitle('Loss', fontsize=20)
plt.legend(loc='upper left')
plt.show()
```

```
In [ ]: fig = plt.figure()
plt.plot(hist.history['accuracy'], color='teal', label='accuracy')
plt.plot(hist.history['val_accuracy'], color='orange', label='val_accuracy')
fig.suptitle('Accuracy', fontsize=20)
plt.legend(loc='upper left')
plt.show()
```

9. Evaluate

```
In [ ]: from tensorflow.keras.metrics import Precision, Recall, BinaryAccuracy
```

```
In [ ]: pre = Precision()
re = Recall()
acc = BinaryAccuracy()
```

```
In [ ]: for batch in test.as_numpy_iterator():
    X, y = batch
    yhat = model.predict(X)
    pre.update_state(y, yhat)
    re.update_state(y, yhat)
    acc.update_state(y, yhat)
```

```
In [ ]: print(pre.result(), re.result(), acc.result())
```

10. Test

```
In [ ]: resize = tf.image.resize(img, (256,256))
plt.imshow(resize.numpy().astype(int))
plt.show()
```

```
In [ ]: yhat = model.predict(np.expand_dims(resize/255, 0))
```

```
In [ ]: yhat
```

```
In [ ]: if yhat > 0.5:
    print(f'Predicted class is Sad')
else:
    print(f'Predicted class is Happy')
```

11. Save the Model

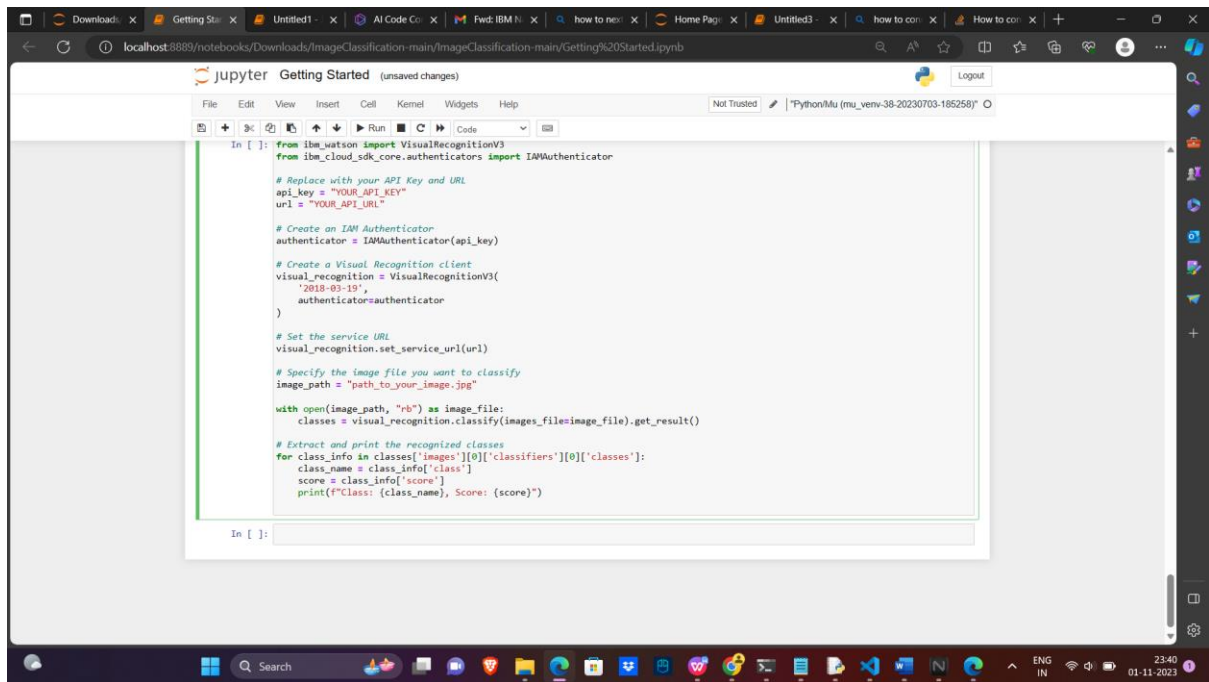
```
In [ ]: from tensorflow.keras.models import load_model
```

```
In [ ]: model.save(os.path.join('models', 'imageclassifier.h5'))
```

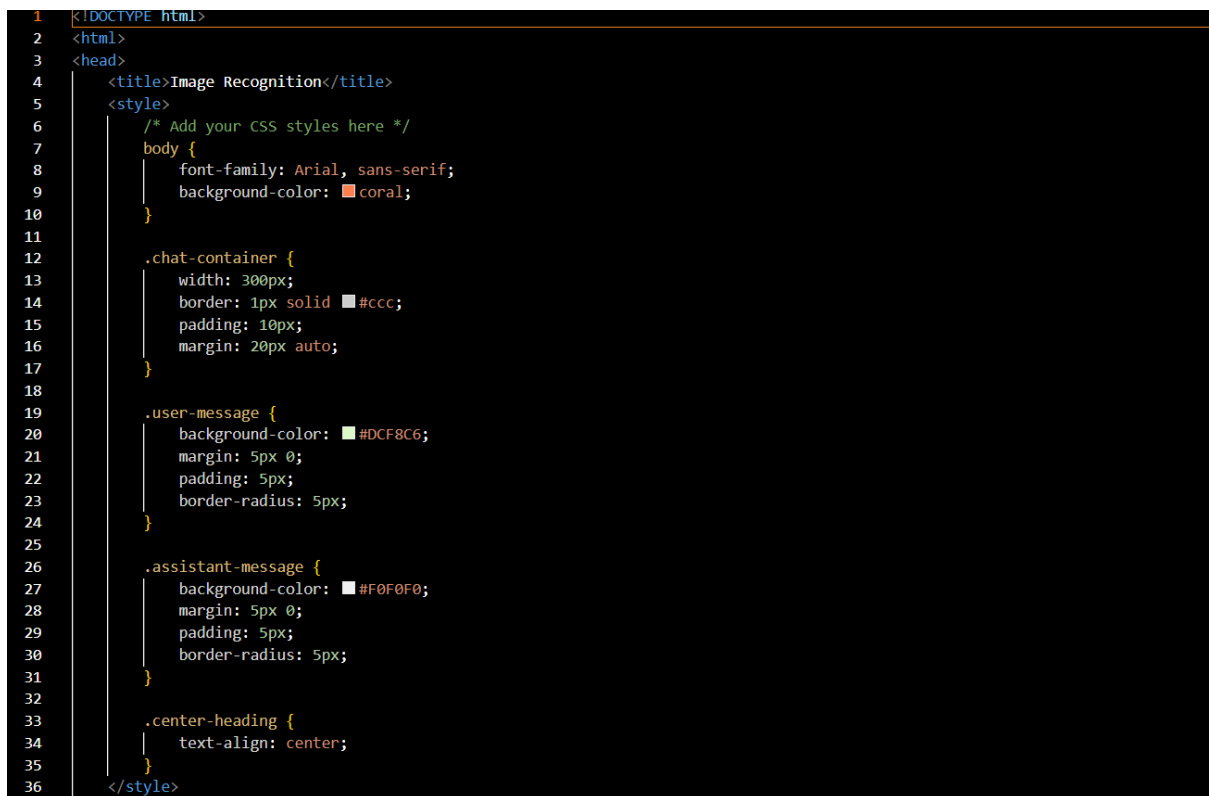
```
In [ ]: new_model = load_model('imageclassifier.h5')
```

```
In [ ]: new_model.predict(np.expand_dims(resize/255, 0))
```

```
In [ ]:
```



CODE FOR WEBSITE LAYOUT:



```

38 </head>
39 <body>
40   <h1 class="center-heading">Image Recognition with AI Captions</h1>
41   <h2 class="center-heading">
42     <input type="file" accept="image/*" id="imageUpload" onchange="processImage()">
43     <img id="uploadedImage" style="max-width: 300px; display: none;">
44   </h2>
45   <h2 class="center-heading">Classified Objects:</h2>
46   <ul id="classifications"></ul>
47
48   <h2 class="center-heading">Generated Caption:</h2>
49   <p id="caption"></p>
50
51   <script>
52     function processImage() {
53       const imageUpload = document.getElementById('imageUpload');
54       const uploadedImage = document.getElementById('uploadedImage');
55       const classifications = document.getElementById('classifications');
56       const caption = document.getElementById('caption');
57
58       const file = imageUpload.files[0];
59       if (!file) {
60         return;
61       }
62
63       // Display the uploaded image
64       const imageURL = URL.createObjectURL(file);
65       uploadedImage.src = imageURL;
66       uploadedImage.style.display = 'block';
67
68       // Perform image classification using IBM Cloud Visual Recognition
69       // You'll need to use the IBM Visual Recognition API here
70
71       // Generate captions for recognized objects
72       // You'll need to use your chosen natural language generation library
73
74       // Update the web page with the classifications and caption
75
76       // Display the uploaded image
77       const imageURL = URL.createObjectURL(file);
78       uploadedImage.src = imageURL;
79       uploadedImage.style.display = 'block';
80
81       // Perform image classification using IBM Cloud Visual Recognition
82       // You'll need to use the IBM Visual Recognition API here
83
84       // Generate captions for recognized objects
85       // You'll need to use your chosen natural language generation library
86
87       // Update the web page with the classifications and caption
88       classifications.innerHTML = "<li>Object 1: DOG</li><li>Object 2: Confidence 0.85</li>";
89       caption.innerHTML = "A generated caption for the recognized objects.";
90
91       // Handle API calls and caption generation here
92     }
93   </script>
94 </body>
95 </html>

```

HOME PAGE:

Image Recognition with AI Captions

Choose File

No file chosen

Classified Objects:


Generated Caption:

OUTPUT PAGE:

Image Recognition with AI Captions

Choose File

download.jpeg



Classified Objects:

Generated Caption:

- Object 1: DOG
- Object 2: Confidence 0.85

A generated caption for the recognized objects.